

AI2000/AI5000/CS5590  
Foundations of Machine Learning

# Course Introduction

Aug 2021

Vineeth N Balasubramanian



# A few (not so, yet) recent quotes

- “A breakthrough in machine learning would be worth ten Microsofts” (Bill Gates, Chairman, Microsoft)
- “Machine learning is the next Internet” (Tony Tether, Director, DARPA)
- “Machine learning is today’s discontinuity” (Jerry Yang, ex-CEO, Yahoo)
- “Machine learning is today’s electricity” (Andrew Ng)

# What is Machine Learning?




# What is Machine Learning?

- Making predictions or decisions from data
- “Programming computers to optimize a performance criterion using example data or past experience” (Ethem Alpaydin, Machine Learning, 2010)
- “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E.” (Tom Mitchell, Machine Learning, 1997)
- “Learning general models from a data of particular examples”
- “Build a model that is *a good and useful approximation* to the data.”

# Today

## Traditional Programming



## Machine Learning



Source: Domingos


# Related Terms

Machine Learning, Data Mining, Knowledge Discovery,  
Artificial Intelligence, Statistical Learning, Pattern Recognition,  
Computational Learning



# When is Machine Learning Used?

- Human expertise does not exist
  - E.g. navigating on Mars
- Humans are unable to explain their expertise
  - E.g. speech recognition
- Solution changes in time
  - E.g. routing on a computer network
- Solution needs to be adapted to particular cases
  - E.g. user biometrics
- Data is cheap and abundant; knowledge is expensive and scarce



# Applications of Machine Learning

From: cheapsales@buystufffromme.com  
To: ang@cs.stanford.edu  
Subject: Buy now!

Deal of the week! Buy now!  
Rolex w4tchs - \$100  
Med1cine (any kind) - \$50  
Also low cost M0rgages available.


Spam

From: Alfred Ng  
To: ang@cs.stanford.edu  
Subject: Christmas dates?

Hey Andrew,  
Was talking to Mom about plans for Xmas. When do you get off work. Meet Dec 22?  
Alf

Non-spam


# Applications of Machine Learning



# Applications of Machine Learning

CCAGCTGCATCACAGGAGGCCAGCGAGCAGGTCTGTTCCAAGGGCCTCGAGCCAGTCTG EI  
GAGGTGAAGGACGTCTCCCCAGGAGCCGGTGAGAAGCGCAGTCGGGGCACGGGGATG EI  
TAAATTCTCTGTTAACACCTTCAGACTTATGTGTATGAAGGAGTAGAAGCCAAA IE  
AAACTAAAGAATTATTCTTACATTCAGTTCTGATCATGAAAACGCCAACAAAA IE  
AAAGCAGATCAGCTGTATAAACAGAAAATTATTCGTGGTTCTGTCACTGTGTGATGGT N  
TTGCCCTCAGCATCACCATGAACGGAGAGGCCATGCCCTGCGCTGAGGGCTGCCAGGCCA N

# Applications of Machine Learning



News:



Words like **Jackson** and **antitrust** are more likely in the stories preceding the plunge.

$P(\text{shares}) = 0.074$   
 $P(\text{antitrust}) = 0.009$   
 $P(\text{judge}) = 0.006$   
 $P(\text{trading}) = 0.032$   
 $P(\text{against}) = 0.025$   
 $P(\text{Jackson}) = 0.001$

Software giant Microsoft saw its shares dip a few percentage points this morning after U.S. District Judge Thomas Penfield Jackson issued his "findings of fact" in the government's ongoing antitrust case against the Seattle wealth-creation machine...



$P(\text{shares} \mid \text{MSFT}\downarrow) = 0.071$   
 $P(\text{antitrust} \mid \text{MSFT}\downarrow) = 0.044$   
 $P(\text{judge} \mid \text{MSFT}\downarrow) = 0.039$   
 $P(\text{trading} \mid \text{MSFT}\downarrow) = 0.029$   
 $P(\text{against} \mid \text{MSFT}\downarrow) = 0.027$   
 $P(\text{Jackson} \mid \text{MSFT}\downarrow) = 0.025$

$$P(\text{MSFT}\downarrow \mid \text{Jackson}) = P(\text{Jackson} \mid \text{MSFT}\downarrow) P(\text{MSFT}\downarrow) / P(\text{Jackson})$$

# Applications of Machine Learning

NETFLIX

English ▾

Sign In

Unlimited movies, TV shows and more.

Watch anywhere. Cancel anytime.

Ready to watch? Enter your email to create or restart your membership.

Email address

Get Started >

# More ML Applications

- Science (Astronomy, neuroscience, medical imaging, bio-informatics)
- Environment (energy, climate, weather, resources)
- Retail (Intelligent stock control, demographic store placement)
- Manufacturing (Intelligent control, automated monitoring, detection methods)
- Security (Intelligent smoke alarms, fraud detection)
- Marketing (promotions, ...)
- Management (Scheduling, timetabling)
- Finance (credit scoring, risk analysis...)
- Web data (information retrieval, information extraction, ...)

# More Recent ML Applications

- AlphaGo!
- Automating Employee Access Control
- Identifying whales in ocean based on audio recordings
- Predict wait times for patients in emergency rooms
- Extract heart failure diagnosis criteria from free-text physician notes
- Predicting hospital readmissions
- Is (s)he a psychopath?

Source: <http://www.forbes.com/sites/85broads/2014/01/06/six-novel-machine-learning-applications/#6b6f9a9e67bf>




# When are ML algorithms not needed?

- When the relationships between all system variables (input, output, and hidden) is completely understood!
- This is NOT the case for almost any real system!


# Overview of ML

- Supervised learning
  - Predict an output  $y$  when given an input  $x$
  - For categorical  $y$  : classification.
  - For real-valued  $y$  : regression.
- Unsupervised learning
  - Create an internal representation of the input, e.g. clustering, dimensionality
  - This is important in machine learning as getting labels is often difficult and expensive
- Other settings of ML
  - Reinforcement learning (learning from “rewards”)
  - Semi-supervised learning (combines supervised + unsupervised)
  - Active learning, Transfer learning, Structured prediction


# Classification (Supervised Learning)




# Classification (Supervised Learning)



# Regression (Supervised Learning)




# Clustering (Unsupervised Learning)



# Dimensionality Reduction (Unsupervised Learning)

- Large sample size is required for high-dimensional data
- Query accuracy and efficiency degrade rapidly as the dimension increases
- Strategies
  - Feature reduction
  - Feature selection
  - Manifold learning
  - Kernel learning



# Other Settings: Ranking (Supervised Learning)

Given a query and a set of web pages, rank them according to relevance


The screenshot shows a Google search results page for the query "machine learning". The results are as follows:

- Machine learning - Wikipedia, the free encyclopedia**  
en.wikipedia.org/wiki/Machine\_learning ▾  
Machine learning, a branch of artificial intelligence, concerns the construction and study of systems that can learn from data. For example, a machine learning ...  
Artificial intelligence - Supervised learning - List of machine learning ... - Weka  
Franck Demontcourt +1'd this
- CS 229: Machine Learning**  
cs229.stanford.edu/ ▾  
Check out this year's awesome projects at Fall 2012 Projects. Come check out the cool new projects during the CS229 Poster Session this Thursday December ...  
You've visited this page 2 times. Last visit: 8/14/13
- Machine Learning | Coursera**  
https://www.coursera.org/course/ml ▾  
Machine learning is the science of getting computers to act without being explicitly programmed. In the past decade, machine learning has given us self-driving ...  
Franck Demontcourt and 3 other people +1'd this
- Machine Learning Department - Carnegie Mellon University**  
www.ml.cmu.edu/ ▾  
Large group with projects in robot learning, data mining for manufacturing and in multimedia databases, causal inference, and disclosure limitation.
- Machine Learning - MIT OpenCourseWare**  
ocw.mit.edu › Courses › Electrical Engineering and Computer Science ▾  
6.867 is an introductory course on machine learning which gives an overview of many concepts, techniques, and algorithms in machine learning, beginning with ...

- Other applications
  - User preference, e.g. Netflix “My List” -- movie queue ranking
  - Flight search (search in general)
  - ...

# Other Settings: Reinforcement Learning

- Learning a policy: A **sequence** of outputs
- No supervised output but delayed reward
  - E.g. Game playing
  - E.g. Robot in a maze
- Multiple agents, partial observability, ...
- Example (Simple Demo):
  - <https://www.youtube.com/watch?v=DCjbk4m1G6I>




# ML Problems

	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>
<i>Discrete</i>	classification or categorization	clustering
<i>Continuous</i>	regression	dimensionality reduction

# ML in Practice


- Understanding domain, prior knowledge, and goals
- Data integration, selection, cleaning, pre-processing, etc.
- Learning models
- Interpreting results
- Consolidating and deploying discovered knowledge
- Loop

# Training and Testing ML Models



# Training and Testing ML Models

- Training is the process of making the system able to learn.
- **No free lunch** rule:
  - Training set and testing set may not come from the same distribution
  - Need to make some assumptions or bias



# Types of Models

- Inductive vs Transductive Learning
- Online vs Offline Learning
- Generative vs Discriminative Models
- Parametric vs Non-Parametric Models

# ML Datasets

- UCI Repository:  
<http://www.ics.uci.edu/~mlearn/MLRepository.html>
- Statlib: <http://lib.stat.cmu.edu/>
- Kaggle
- Many more...

# ML Resources

- MOOCs
  - Coursera, EdX, Udacity
- Conferences/Journals
  - JMLR, Machine Learning, IEEE Transactions on Neural Networks and Learning Systems, IEEE Transactions on Pattern Analysis and Machine Intelligence, Annals of Statistics
  - ICML, NIPS, KDD, IJCAI, AAAI, ICDM

# Mathematical Basis

- Functions, Logarithms and Exponentials
- Vectors, Dot Products, Orthogonality
- Matrices, Matrix Operations, Linear Transformations, Eigendecomposition
- Calculus, Differentiation, Integration
- Probability and Statistics
- Functional Analysis, Hilbert Spaces

# Course Details

- Please see [course web link](#)

# Topics not to be covered (likely)

- Deep Learning (not in its “depth”, at least)
- Probabilistic Graphical Models/Bayesian Networks (not in depth)
- Reinforcement Learning

# Programming

- Python
- Libraries
  - Numpy, Scipy – numerical/scientific computing, linear algebra
  - Matplotlib – for plotting
  - Scikitlearn – for machine learning

# Foundational Reading/Follow-up

- Math
  - Part 1 of Deep Learning book: <http://www.deeplearningbook.org/>
  - Essence of linear algebra: <http://youtu.be/kjBOesZCoqc>
  - Essence of calculus: <https://goo.gl/Hnk1jA>
- Programming
  - Learn Python
    - <https://try.jupyter.org/>
    - <https://docs.python.org/3/tutorial/>
    - Video Tutorials: <https://www.youtube.com/watch?v=cpPG0bKHYKc>
  - Play with Numpy